

# **Advanced Data Analysis**

DATA 71200

Class 3

# Schedule

5-Jun	<b>Inspecting Data</b>
6-Jun	Async: DataCamp
10-Jun	Evaluation Methods
11-Jun	Async: DataCamp
12-Jun	Supervised Learning (k-Nearest Neighbors, Linear Models, and Naive Bayes Classifiers)
13-Jun	<i>Async: DataCamp Modules Project 1 Due</i>

# Inspecting Data to Gain Insights

## ► Review from last class

- Data size and type
- Summary statistics
- Histograms
- Scatter Matrix

# Representing Data

- ▶ **Continuous versus categorical**
  - One-Hot Encoding
  - Binning
- ▶ **Transformations**
- ▶ **Automatic feature selection**
- ▶ **Utilizing expert knowledge**


# Some Terminology

## ▸ (Linear) Regression

- Continuous predictive model created by estimating a linear relationship between features

## ▸ Logistic Regression

- Predictive model of the probability of a certain class



# Some Terminology

## ► **Regularization**

- Adds an extra term to the cost function
- Can be applied to linear and logistic regression
- Can also be used for feature selection
- Lasso (least absolute shrinkage and selection operator) regression is another form, referred to as L1
- Ridge is a form of regularization, referred to as L2

# Some Terminology

- ▶ **Lasso Regression (L1)**

- reduces the coefficients of the least important variables to zero (removing them completely by the model)

- ▶ **Ridge Regression (L2)**

- addresses *multicollinearity* (linear relationships between parameters) and having more parameters than observations

# Continuous Versus Categorical

- ▶ **(Linear) Regression - predicts continuous values**
- ▶ **Classification - predicts categorical, or discrete, values**
- ▶ **Continuous versus categorical distinction also holds for input features**



# One-Hot Encoding

- ▶ Split the different categories in their own variable
- ▶ E.g., a single variable for color where the values are the strings “blue”, “red”, “yellow” would be encoded as

	Blue	Red	Yellow
Blue	1	0	0
Red	0	1	0
Yellow	0	0	1

← **Variables**

↑  
**Values**

*Categorical data can also be encoded as numbers*

	Categorical Feature	Integer Feature
0	socks	0
1	fox	1
2	socks	2
3	box	1

**1      2      3      4      5      6**

**0      1      2      3      4      5**

```
[ [ 0.  0.  1.  1.  0.  0. ]
  [ 0.  1.  0.  0.  1.  0. ]
  [ 0.  0.  1.  0.  0.  1. ]
  [ 1.  0.  0.  0.  1.  0. ] ]
```

**box   fox   socks   0      1      2**

# In-Class Activity 1

- ▶ **Apply one-hot encoding to the ocean\_proximity value in the California Housing dataset that we looked at last class**
  - Using `pd.dummies` and/or `OneHotEncoder` from `scikitlearn`
  - `housing['ocean_proximity'].values.reshape(-1,1)`

# Binning

- ▶ **Discretizing continuous data into numerical bins can be useful when small differences in value are not significant**
- ▶ **E.g., for numerical grade data (out of 100), it may be more useful to give a model how many scores fall into ranges of 5 rather than the continuous data**

82	83	92	93	72	73	87	86	99	97	98	51	52	82	81	87	91	92	61	67	
										↓										
50-54	55-59	60-64	65-69	70-74	75-79	80-84	85-89	90-94	95-99											
2	0	1	2	2	0	4	3	4	2											

# In-Class Activity 2

- ▶ **Apply binning to the housing\_median\_age value in the California Housing dataset that we looked at last class**
  - `housing['housing_median_age'].values.reshape(-1, 1)`
  - Plot both the original data and the binned data
- ▶ **Explore binning with other features**

# Transformations

- ▶ **Squaring and cubing is useful for linear regression models**
- ▶ **Logarithms and exponentials are useful for representing your data with a Gaussian distribution, which is useful for mean-based models**

# In-Class Activity 3

- ▶ **Apply the following transformations to `housing_median_age` in the California Housing dataset that we looked at last class**
  - Squaring (`**2`)
  - Cubing (`**3`)
  - `np.log`
  - `np.exp`
- ▶ **Plot histograms and scatter matrices to explore the resultant data (for `**2`, `**3`, and `np.log`)**

# Automatic Feature Selection

- ▶ **Regularization can be used to assess the relative importance of features in the performance of a model**
  - Although this can't tell you anything about features you don't include
- ▶ **Recursive feature elimination (RFE) starts with all features and removes the poorly performing ones**
- ▶ **You can also start with one feature and build up a model**



# Utilizing Expert Knowledge

- ▶ **Domain knowledge can be useful for recognizing patterns in data that may be beneficial or detrimental to the model**
- ▶ **This can inform decisions about which features to include and how to represent them**

# Reading

- Ch 4: "Representing Data/Engineering Features" in Guido, Sarah and Andreas C. Muller. (2016). Introduction to Machine Learning with Python, O'Reilly Media, Inc. 213–55.

# DataCamp for next class

- ▶ *Preprocessing for Machine Learning in Python*

# Project 1

- ▶ **Due June 13**
- ▶ **Keep exploring potential datasets**
  - [kaggle.com](https://www.kaggle.com)
  - [archive.ics.uci.edu/ml/datasets.php](https://archive.ics.uci.edu/ml/datasets.php)
  - [libguides.nypl.org/eresources](https://libguides.nypl.org/eresources)
  - [opendata.cityofnewyork.us/data/](https://opendata.cityofnewyork.us/data/)
- ▶ **The data set will need to be labeled as you are going to use it for both supervised and unsupervised learning tasks**
- ▶ ***We will go over using IMPUTER to address missing values on Monday***

# Project 1

Curation and cleaning of a labeled data set that you will use for the supervised and unsupervised learning tasks in project 2 and 3. The dataset can be built from existing data and should be stored in your GitHub repository.

To **submit** the assignment submit a link to your project Jupyter notebook on Blackboard

The **goal** for this assignment is for you to create a usable dataset from an open-source data collection. You will use your curated dataset for a supervised classification task in Project 2 and an unsupervised learning task in Project 3.

**Step 1: Find and download a dataset.** Here are some potential places to look

- Amazon's AWS datasets: <https://aws.amazon.com/opendata/public-datasets/>
- Data Portals: <http://dataportals.org/>
- Kaggle datasets: <http://kaggle.com>
- NYPL digitizations: <http://libguides.nypl.org/eresources>
- NYC Open Data: <http://opendata.cityofnewyork.us/data/>
- Open Data Monitor: <http://opendatamonitor.eu/>
- Quandl: <http://quandl.com/>
- UC Irvine Machine Learning Repository: <https://archive.ics.uci.edu/ml/index.php>

*Some guidelines regarding dataset selection:*

- Since this data is going to be used for supervised learning, one of its features should be a value that you can predict (e.g., if you chose a real estate dataset and are interested in housing prices, then housing prices for each sample should be available in the dataset)
- Your dataset should be large enough to perform supervised and unsupervised learning on. A useful rule of thumb is that you should have at least 10 samples per dimension or parameter/feature you are fitting.

**Step 2: Divide into a training set and a testing set.** In a Jupyter notebook, use scikitlearn to divide your data into training and testing sets. Make sure that the testing and training sets are balanced in terms of target classes

```
1 from sklearn.model_selection import train_test_split
2
3 # split data and labels into a training and a test set
4 X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=0, stratify=None)
5
6 #a balanced split, percentage of samples for each class, can be obtained with StratifiedShuffleSplit
7 #note however that the housing dataset is not a good candidate for this approach
8 from sklearn.model_selection import StratifiedShuffleSplit
9 split = StratifiedShuffleSplit(n_splits=1, test_size=0.2, random_state=42)
10 for train_index, test_index in split.split(X, y):
11     X_train = X[train_index]
12     X_test = X[test_index]
13     y_train = y[train_index]
14     y_test = y[test_index]
```

**data71200class4.ipynb**

**Step 3: Explore your training set.** In a Jupyter notebook, import your data into a Pandas data frame and use the following pandas functions to explore your data

- DataFrame.info()
- DataFrame.describe()

**Step 4: Data cleaning.** Address any missing values in your training set. Include the code in your Jupyter notebook and create a second, cleaned, version of your dataset. Then apply the same procedure to your test set (if you are putting in replacement values use IMPUTER in scikitlearn).

- Recall from the *Hands on Machine Learning* book (p. 60) that some options are
  - “Get rid of the corresponding samples.”
  - “Get rid of the whole attribute (column).”
  - “Set the values to some value (zero, the mean, the median, etc.).”

## Dealing with missing value

data71200class4.ipynb

```
1 X = housing.drop(['median_house_value', 'ocean_proximity'], axis=1)
```

```
1 from sklearn.impute import SimpleImputer
2 imp_mean = SimpleImputer(missing_values=np.nan, strategy='mean')
3 imp_mean.fit(X)
4 SimpleImputer()
5 X = imp_mean.transform(X)
```

```
1 # instantiate a model and fit it to the training set
2 linreg = LinearRegression().fit(X_train, y_train)
```

```
1 # instantiate a model and fit it to the training set
2 linreg = LinearRegression().fit(X_train, y_train)
```

Test set score: 0.63

```
1 # evaluate the model on the test set
2 print("Test set score: {:.2f}".format(linreg.score(X_test, y_test)))
```



## If you want to see which feature has the missing element

```
1 # you can check each one individually with the following code
2 housing['longitude'].isnull().values.any()
```

False

```
1 housing['latitude'].isnull().values.any()
```

False

```
1 housing['housing_median_age'].isnull().values.any()
```

False

```
1 housing['total_rooms'].isnull().values.any()
```

False

```
1 housing['total_bedrooms'].isnull().values.any()
```

True

```
1 # although this isn't necessary for running Imputer it may be useful to know where exactly the data is missing
2 # you can also check how many elements are empty
3 housing['total_bedrooms'].isnull().values.sum()
```

207

```
1 # or you can generally check if you have any empty elements in your dataframe
2 #(and thus whether you need to run Imputer)
3 housing.isnull().values.any()
```

True

**Step 5: Visualize the data in your training set.** At a minimum, use the following pandas functions to visualize the data in your Jupyter notebook.

- `DataFrame.hist`
- `plotting.scatter_matrix()`

**data71200class2lab.ipynb**

**Step 6: Apply transformations to your data.** In your Jupyter notebook apply, squaring, cubing, logarithmic, and exponential transformations to two features in your dataset. Plot the histograms and scatter matrices of the resultant data.

**data71200class3lab.ipynb**