

Advanced Data Analysis

DATA 71200

Class 4

Schedule

8-Jun Evaluation Methods

9-Jun Async: DataCamp

13-Jun Supervised Learning
(k-Nearest Neighbors and Linear Models)

14-Jun Async: DataCamp

Schedule

10-Jun	Evaluation Methods
11-Jun	Async: DataCamp
12-Jun	Supervised Learning (k-Nearest Neighbors, Linear Models, and Naive Bayes Classifiers)
13-Jun	<i>Async: DataCamp Modules</i> <i>Project 1 Due</i>
17-Jun	Supervised Learning (Decision Trees, Support Vector Machines and Uncertainty estimates from Classifiers)
18-Jun	<i>Async: DataCamp Modules</i>

Project 1

Curation and cleaning of a labeled data set that you will use for the supervised and unsupervised learning tasks in project 2 and 3. The dataset can be built from existing data and should be stored in your GitHub repository.

To **submit** the assignment submit a link to your project Jupyter notebook on Blackboard

The **goal** for this assignment is for you to create a usable dataset from an open-source data collection. You will use your curated dataset for a supervised classification task in Project 2 and an unsupervised learning task in Project 3.

Step 1: Find and download a dataset. Here are some potential places to look

- Amazon's AWS datasets: <https://aws.amazon.com/opendata/public-datasets/>
- Data Portals: <http://dataportals.org/>
- Kaggle datasets: <http://kaggle.com>
- NYPL digitizations: <http://libguides.nypl.org/eresources>
- NYC Open Data: <http://opendata.cityofnewyork.us/data/>
- Open Data Monitor: <http://opendatamonitor.eu/>
- QuandDL: <http://quandl.com/>
- UC Irvine Machine Learning Repository: <https://archive.ics.uci.edu/ml/index.php>

Some guidelines regarding dataset selection:

- Since this data is going to be used for supervised learning, one of its features should be a value that you can predict (e.g., if you chose a real estate dataset and are interested in housing prices, then housing prices for each sample should be available in the dataset)
- Your dataset should be large enough to perform supervised and unsupervised learning on. A useful rule of thumb is that you should have at least 10 samples per dimension or parameter/feature you are fitting.

Step 2: Divide into a training set and a testing set. In a Jupyter notebook, use scikitlearn to divide your data into training and testing sets. Make sure that the testing and training sets are balanced in terms of target classes

```
1 from sklearn.model_selection import train_test_split
2
3 # split data and labels into a training and a test set
4 X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=0, stratify=None)
5
6 #a balanced split, percentage of samples for each class, can be obtained with StratifiedShuffleSplit
7 #note however that the housing dataset is not a good candidate for this approach
8 from sklearn.model_selection import StratifiedShuffleSplit
9 split = StratifiedShuffleSplit(n_splits=1, test_size=0.2, random_state=42)
10 for train_index, test_index in split.split(X, y):
11     X_train = X[train_index]
12     X_test = X[test_index]
13     y_train = y[train_index]
14     y_test = y[test_index]
```

data71200class4.ipynb

Step 3: Explore your training set. In a Jupyter notebook, import your data into a Pandas data frame and use the following pandas functions to explore your data

- DataFrame.info()
- DataFrame.describe()

Step 4: Data cleaning. Address any missing values in your training set. Include the code in your Jupyter notebook and create a second, cleaned, version of your dataset. Then apply the same procedure to your test set (if you are putting in replacement values use IMPUTER in scikitlearn).

- Recall from the *Hands on Machine Learning* book (p. 60) that some options are
 - “Get rid of the corresponding samples.”
 - “Get rid of the whole attribute (column).”
 - “Set the values to some value (zero, the mean, the median, etc.).”

Dealing with missing value

data71200class4.ipynb

```
1 X = housing.drop(['median_house_value', 'ocean_proximity'], axis=1)
```

```
1 from sklearn.impute import SimpleImputer
2 imp_mean = SimpleImputer(missing_values=np.nan, strategy='mean')
3 imp_mean.fit(X)
4 SimpleImputer()
5 X = imp_mean.transform(X)
```

```
1 # instantiate a model and fit it to the training set
2 linreg = LinearRegression().fit(X_train, y_train)
```

```
1 # instantiate a model and fit it to the training set
2 linreg = LinearRegression().fit(X_train, y_train)
```

Test set score: 0.63

```
1 # evaluate the model on the test set
2 print("Test set score: {:.2f}".format(linreg.score(X_test, y_test)))
```

If you want to see which feature has the missing element

```
1 # you can check each one individually with the following code  
2 housing['longitude'].isnull().values.any()
```

False

```
1 housing['latitude'].isnull().values.any()
```

False

```
1 housing['housing_median_age'].isnull().values.any()
```

False

```
1 housing['total_rooms'].isnull().values.any()
```

False

```
1 housing['total_bedrooms'].isnull().values.any()
```

True

```
1 # although this isn't necessary for running Imputer it may be useful to know where exactly the data is missing  
2 # you can also check how many elements are empty  
3 housing['total_bedrooms'].isnull().values.sum()
```

207

```
1 # or you can generally check if you have any empty elements in your dataframe  
2 #(and thus whether you need to run Imputer)  
3 housing.isnull().values.any()
```

True

data71200class4.ipynb

Step 5: Visualize the data in your training set. At a minimum, use the following pandas functions to visualize the data in your Jupyter notebook.

- DataFrame.hist
- plotting.scatter_matrix()

data71200class2lab.ipynb

Step 6: Apply transformations to your data. In your Jupyter notebook apply, squaring, cubing, logarithmic, and exponentials transformations to two features in your dataset. Plot the histograms and scatter matrices of the resultant data.

data71200class3lab.ipynb

Submission Instructions

- ▶ **Either upload your data to your GitHub repo or include code in your notebook to download the data from another source**
- ▶ **Create a single notebook for this assignment that**
 - briefly describes your dataset in a markdown (text) cell
 - contains all of the code for Steps 2-6
 - uses markdown (text) cells to indicate which Step (2–6) the code belongs to

Splitting the Data

- ▶ **Imputers fill in missing values in datasets**
- ▶ **Different types of imputes use different strategies to estimate the missing values**
- ▶ **SimpleImputer in Scikit-learn allows you select either the mean, median, or most frequent values**

**Notebook
data71200class4.ipynb
[12–14]**

Splitting the Data

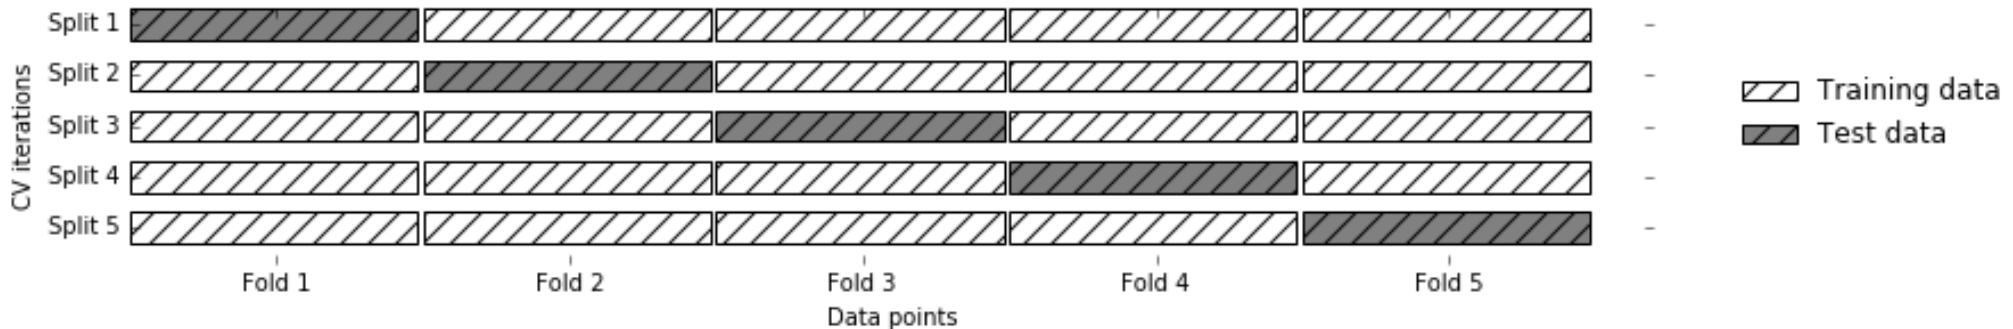


Figure 5-1. Data splitting in five-fold cross-validation

This will create issues if the data is ordered by class since some classes will appear disproportionately (or not at all) in either the testing or training sets

Splitting the Data

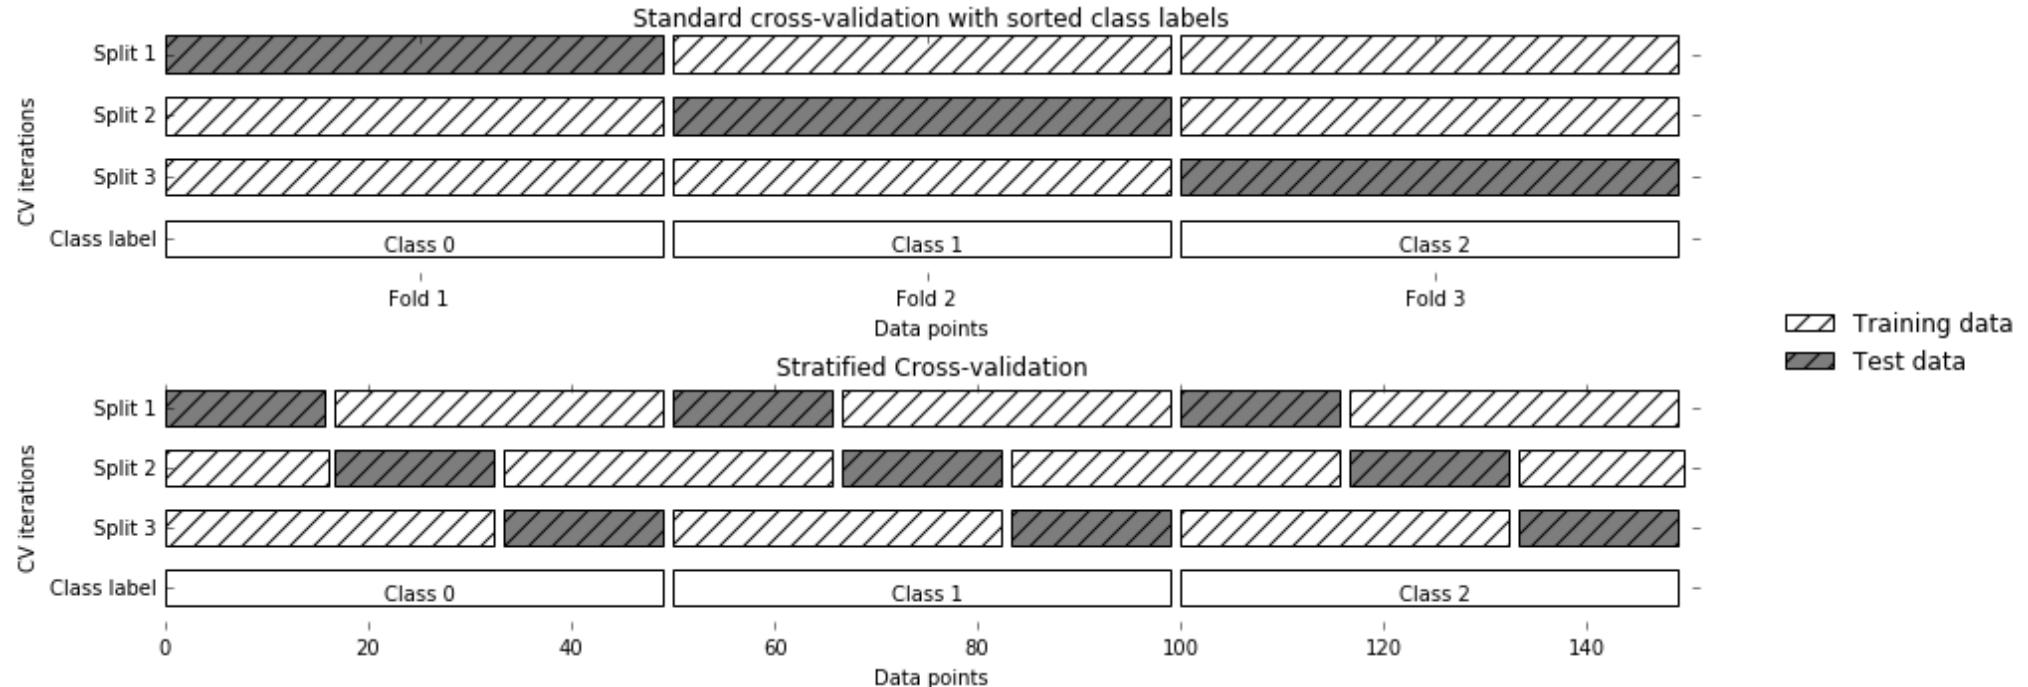


Figure 5-2. Comparison of standard cross-validation and stratified cross-validation when the data is ordered by class label

Stratified Cross-validation explicitly takes data from each class in order to count the issue of over-/under-sampling that can arise in standard cross-validation

Splitting the Data

Small Data Sets

Leave-one out cross validation can be used where each fold is a single sample

Large Data Sets

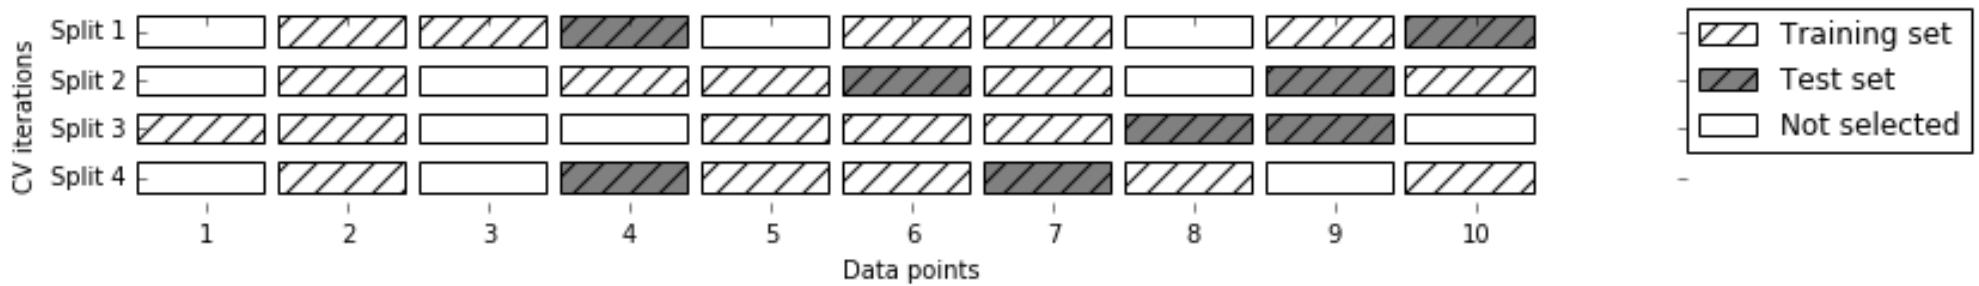


Figure 5-3. *ShuffleSplit* with 10 points, `train_size=5`, `test_size=2`, and `n_splits=4`

Shuffle-Split provides a systematic way to sample from your dataset without using all of the data

Splitting the Data

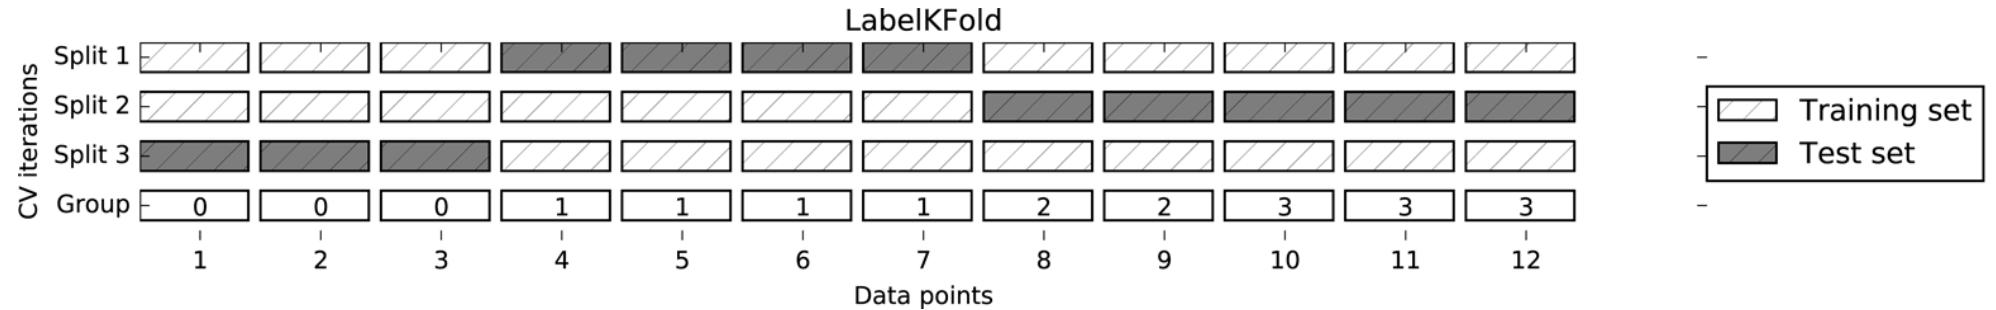


Figure 5-4. Label-dependent splitting with GroupKFold

When groups are highly related (e.g., individual speakers in speech recognition) you may not want to train and test on a single group

GroupKFold allows you select training and testing sets that either include or exclude an entire group

Splitting the Data

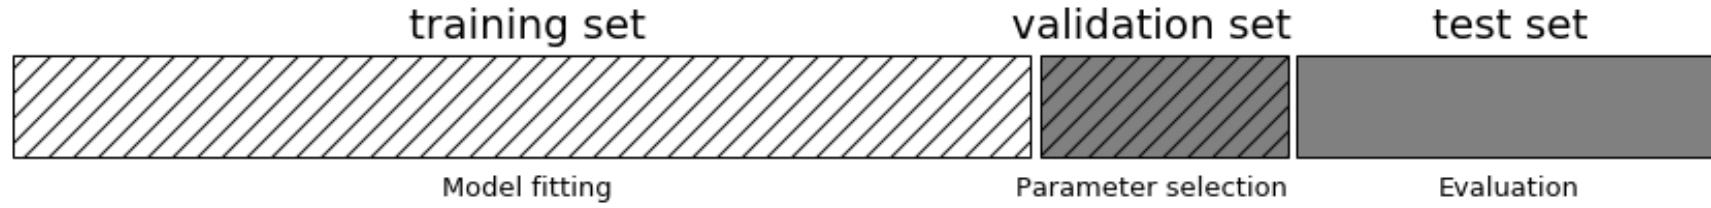


Figure 5-5. A threefold split of data into training set, validation set, and test set

If you are evaluating different models or tuning parameters (e.g., using Grid Search), you will want to have a reserved test set and then split the remaining data into training and validation sets

Grid Search

Grid search can be performed simply (one split) or iteratively (cross-validation) - in the latter the optimal parameter values need to be summarized

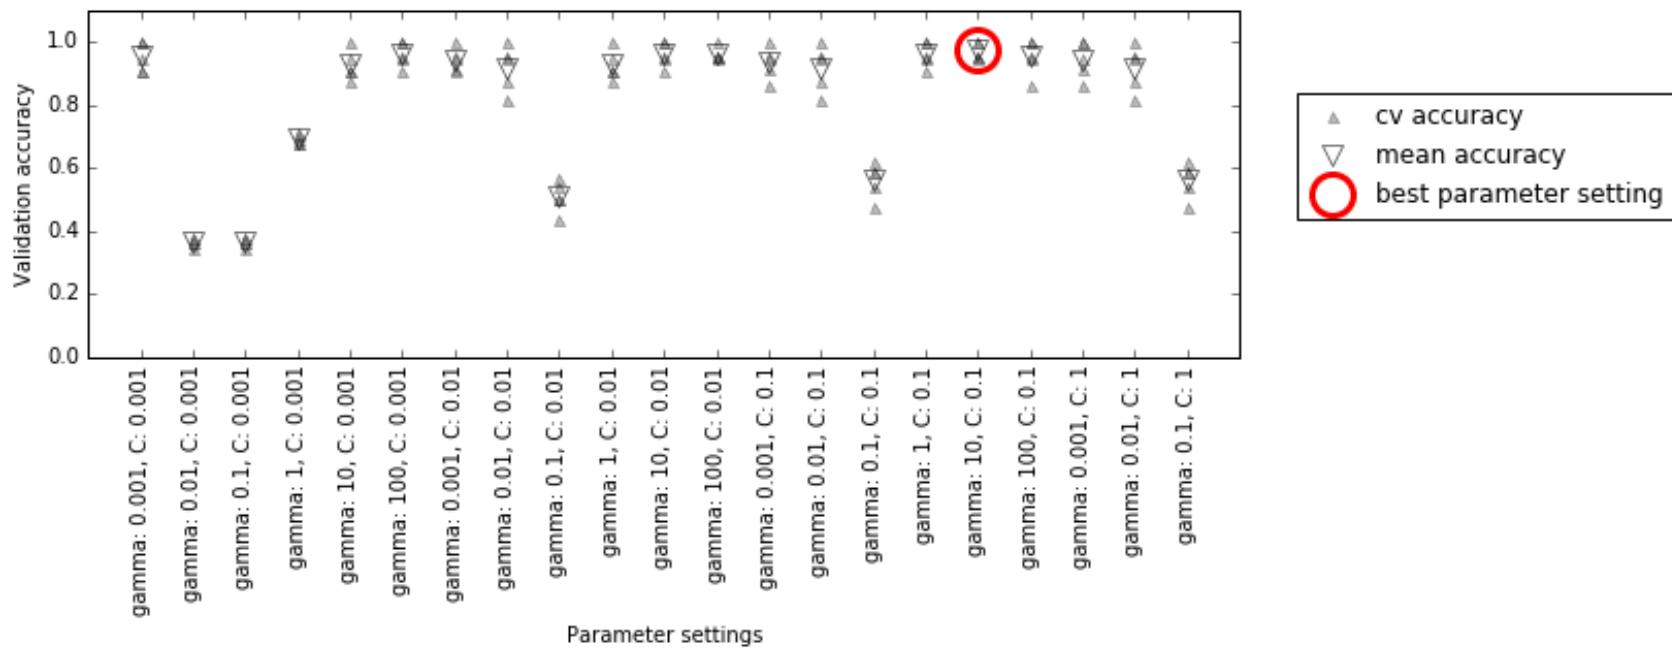


Figure 5-6. Results of grid search with cross-validation

Grid Search

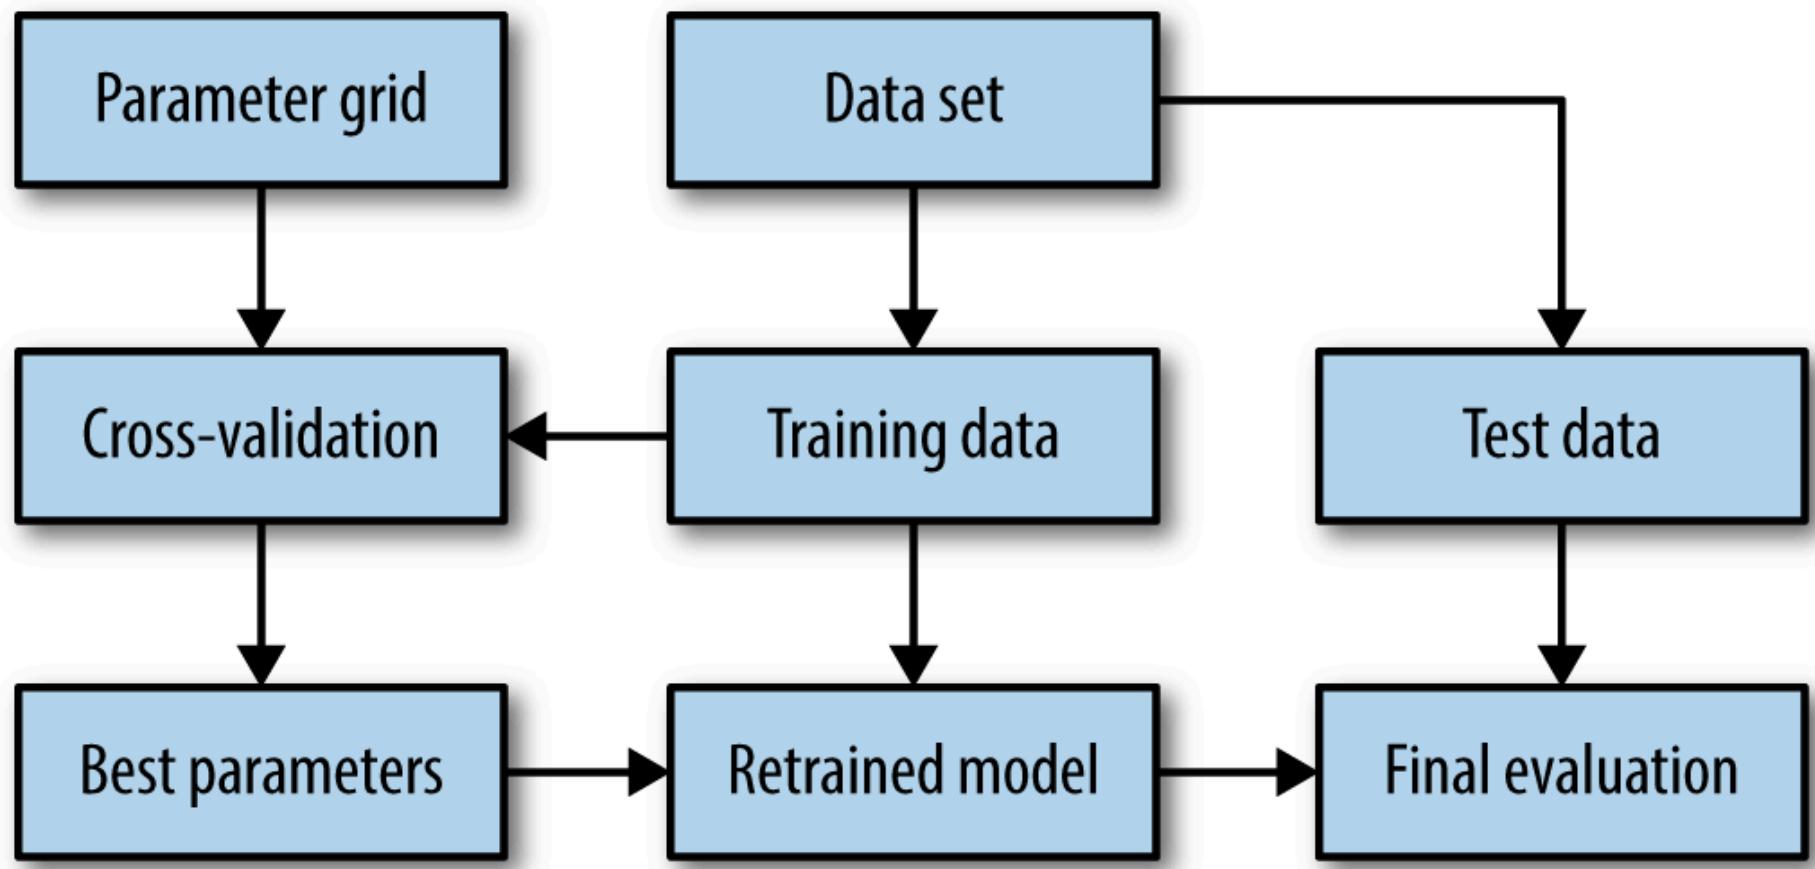


Figure 5-7. Overview of the process of parameter selection and model evaluation with `GridSearchCV`

Grid Search

It is important than the parameter range reached is appropriate

Heat maps can help understand whether the range is appropriate

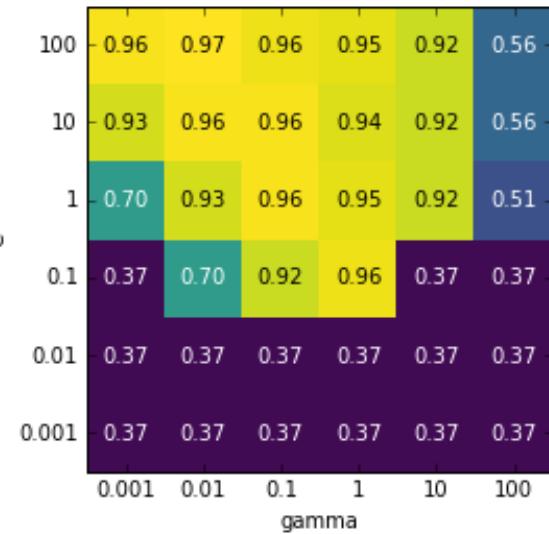


Figure 5-8. Heat map of mean cross-validation score as a function of C and γ

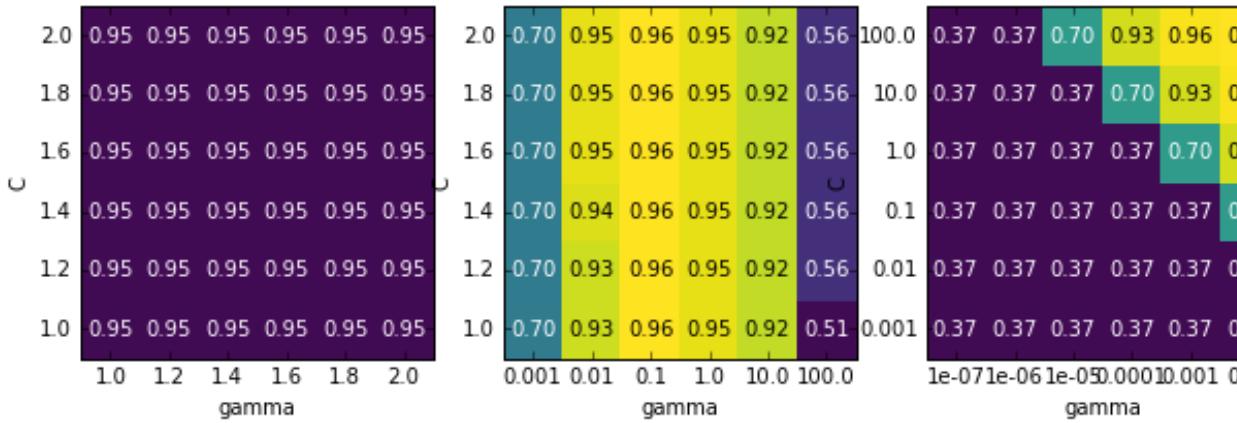
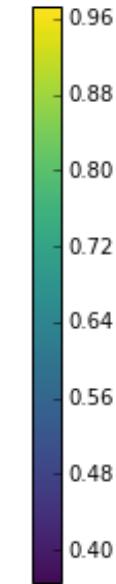


Figure 5-9. Heat map visualizations of misspecified search grids



Missing Values

- ▶ **Cross-validation (within the training set)**

- Splitting the data into folds and iteratively switching up which fold is used for validation (while the remainder are used for training)
- The data can be split up in a variety of ways

**Notebook
data71200class4.ipynb
[12–21]**

Accuracy

$$\text{Accuracy} = \frac{\text{TP} + \text{TN}}{\text{TP} + \text{TN} + \text{FP} + \text{FN}}$$

TP - true positive

TN - true negative

FP - false positive

FN - false negative

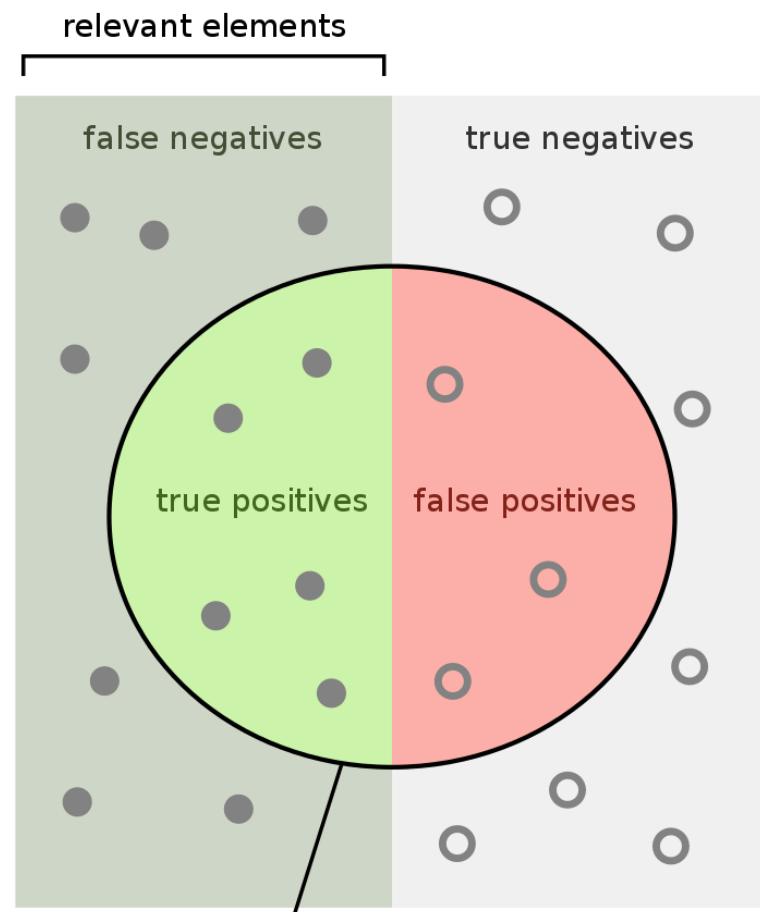
**How many predictions were correct
out of all the predictions made**

Kinds of Errors

Type I: false positive

Type II: false negative

Balance of the classes in the dataset will influence how you interpret performance errors



selected elements

How many relevant items are selected?
e.g. How many sick people are correctly identified as having the condition.

How many negative selected elements are truly negative?
e.g. How many healthy people are identified as not having the condition.

$$\text{Sensitivity} = \frac{\text{True Positives}}{\text{True Positives} + \text{False Negatives}}$$

$$\text{Specificity} = \frac{\text{True Negatives}}{\text{True Negatives} + \text{False Positives}}$$

Confusion Matrices

		True condition	
		Condition positive	Condition negative
Total population			
Predicted condition	Predicted condition positive	True positive	False positive, Type I error
	Predicted condition negative	False negative, Type II error	True negative

Confusion matrices visualize not only how many correct estimates a model made but also what and how it mis-classified

Confusion Matrices

	TN	FP
negative class		
	FN	TP
positive class		
	predicted negative	predicted positive

Confusion matrices visualize not only how many correct estimates a model made but also what and how it mis-classified

Notebook
data71200class4.ipynb
[27]

Precision

$$\text{Precision} = \frac{tp}{tp + fp}$$

tp - true positive

fp - false positive

Positive predictive value

Number of the positive predicted values that are actually positive

Notebook
data71200class4.ipynb
[28]

Recall

$$\text{Recall} = \frac{tp}{tp + fn}$$

tp - true positive

fn - false negative

Sensitivity

Proportion of the actual positives identified

**Notebook
data71200class4.ipynb
[29]**

F-1 Score (F-score, F-measure)

$$F_1 = 2 \cdot \frac{\text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}.$$

Harmonic mean of precision and recall

Summary measurement of the classifier's accuracy

Notebook
data71200class4.ipynb
[30]

Putting it All Together

	True condition			
Total population	Condition positive	Condition negative	Prevalence = $\frac{\sum \text{Condition positive}}{\sum \text{Total population}}$	Accuracy (ACC) = $\frac{\sum \text{True positive} + \sum \text{True negative}}{\sum \text{Total population}}$
Predicted condition	True positive	False positive, Type I error	Positive predictive value (PPV), Precision = $\frac{\sum \text{True positive}}{\sum \text{Predicted condition positive}}$	False discovery rate (FDR) = $\frac{\sum \text{False positive}}{\sum \text{Predicted condition positive}}$
Predicted condition positive	True positive	False positive, Type I error	Positive predictive value (PPV), Precision = $\frac{\sum \text{True positive}}{\sum \text{Predicted condition positive}}$	False discovery rate (FDR) = $\frac{\sum \text{False positive}}{\sum \text{Predicted condition positive}}$
Predicted condition negative	False negative, Type II error	True negative	False omission rate (FOR) = $\frac{\sum \text{False negative}}{\sum \text{Predicted condition negative}}$	Negative predictive value (NPV) = $\frac{\sum \text{True negative}}{\sum \text{Predicted condition negative}}$
	True positive rate (TPR), Recall, Sensitivity, probability of detection, Power $= \frac{\sum \text{True positive}}{\sum \text{Condition positive}}$	False positive rate (FPR), Fall-out, probability of false alarm $= \frac{\sum \text{False positive}}{\sum \text{Condition negative}}$	Positive likelihood ratio (LR+) $= \frac{\text{TPR}}{\text{FPR}}$	Diagnostic odds ratio (DOR) $= \frac{\text{LR+}}{\text{LR-}}$
	False negative rate (FNR), Miss rate $= \frac{\sum \text{False negative}}{\sum \text{Condition positive}}$	Specificity (SPC), Selectivity, True negative rate (TNR) $= \frac{\sum \text{True negative}}{\sum \text{Condition negative}}$	Negative likelihood ratio (LR-) $= \frac{\text{FNR}}{\text{TNR}}$	$F_1 \text{ score} = 2 \cdot \frac{\text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}}$

In-Class Activity

Calculate each of these metrics for the confusion matrix on the right

$$\text{Accuracy} = \frac{\text{TP} + \text{TN}}{\text{TP} + \text{TN} + \text{FP} + \text{FN}}$$

$$\text{Precision} = \frac{tp}{tp + fp}$$

$$\text{Recall} = \frac{tp}{tp + fn}$$

$$F_1 = 2 \cdot \frac{\text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}$$

		Confusion Matrix			
		0	1	2	3
True	Predicted	0	37	0	0
		1	0	39	0
2	0	0	41	3	
3	0	0	1	43	

Figure 5-18. Confusion matrix for the 10-digit classification task

In-Class Activity

	TP	FP	FN
0			
1			
2			
3			
Total			

	0	1	2	3
0	37	0	0	0
1	0	39	0	0
2	0	0	41	3
3	0	0	1	43

True

Predicted

$$\text{Accuracy} = \frac{\text{TP} + \text{TN}}{\text{TP} + \text{TN} + \text{FP} + \text{FN}}$$

$$\text{Precision} = \frac{tp}{tp + fp}$$

$$\text{Recall} = \frac{tp}{tp + fn}$$

$$F_1 = 2 \cdot \frac{\text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}$$

Precision Recall Curves

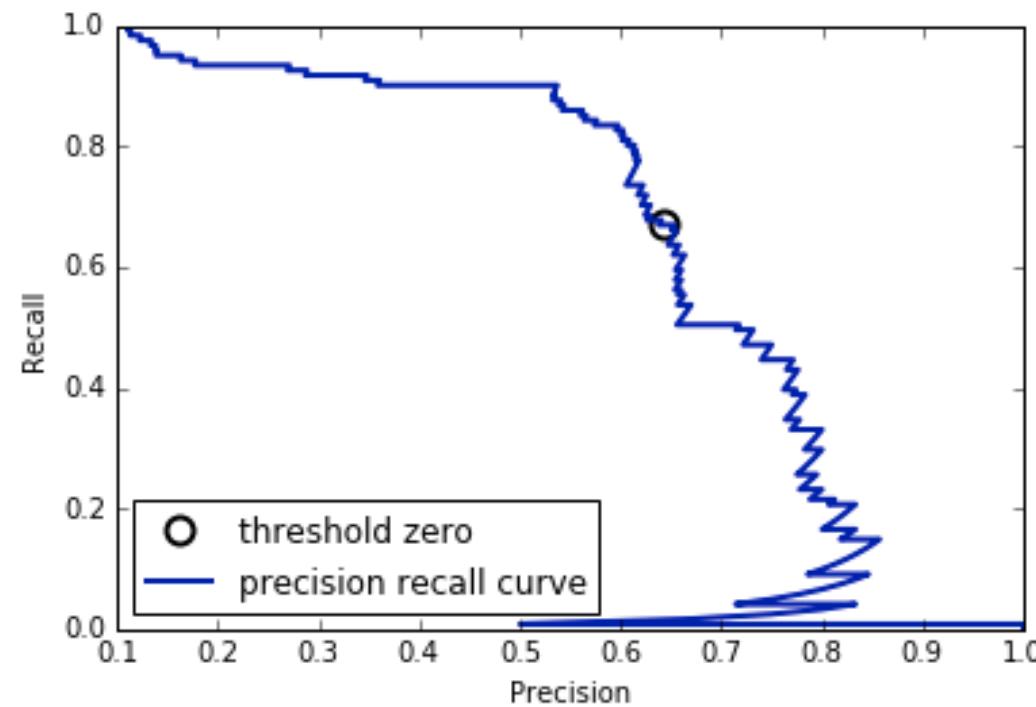


Figure 5-13. Precision recall curve for SVC($\gamma=0.05$)

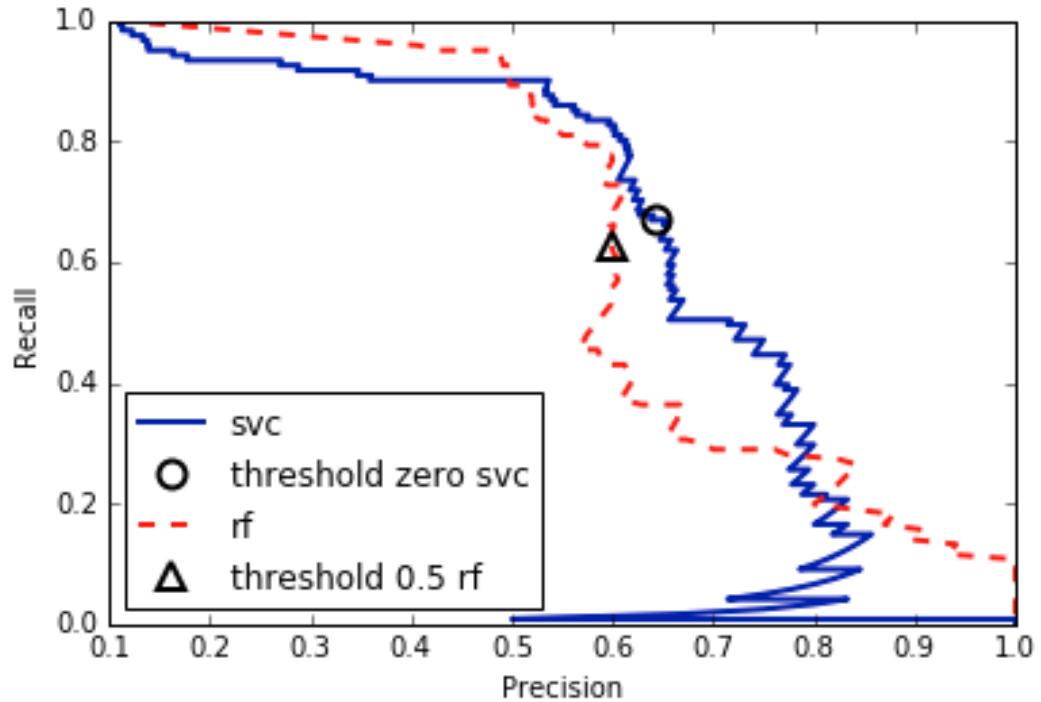


Figure 5-14. Comparing precision recall curves of SVM and random forest

Visualizes trade-off between positive predictive rate and true positive rate (sensitivity)

Area under curve summarizes this information

ROC Curves

$$FPR = \frac{FP}{FP + TN}$$

TP - true positive
TN - true negative
FP - false positive

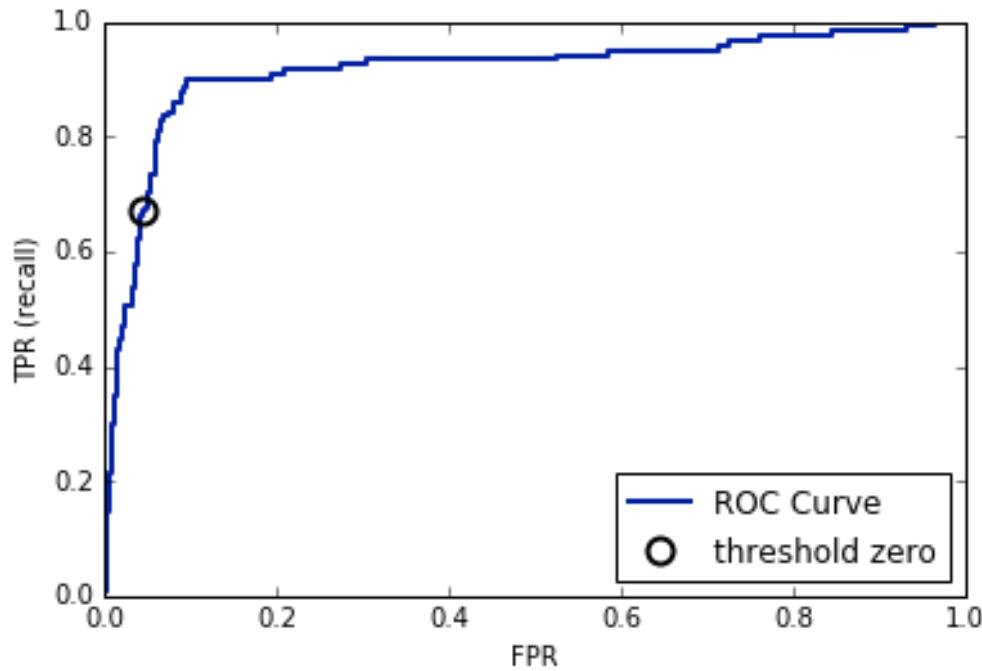


Figure 5-15. ROC curve for SVM

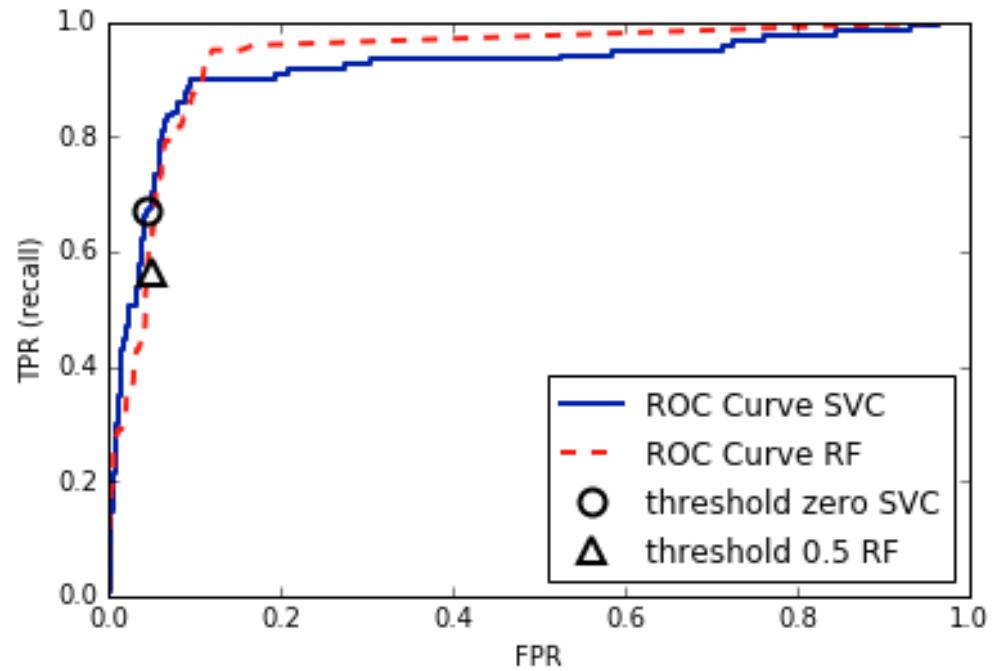


Figure 5-16. Comparing ROC curves for SVM and random forest

Receiver operating characteristics

Visualizes relationship between true positive rate (sensitivity) and false positive rate - the closer to the top right the better

AUC

$$FPR = \frac{FP}{FP + TN}$$

TP - true positive
TN - true negative
FP - false positive

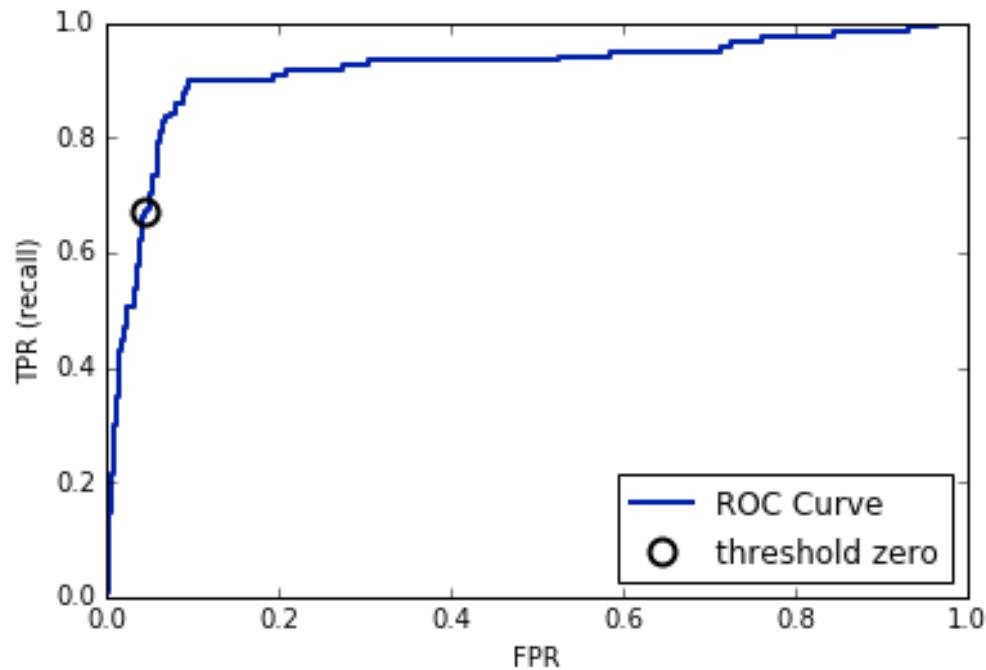


Figure 5-15. ROC curve for SVM

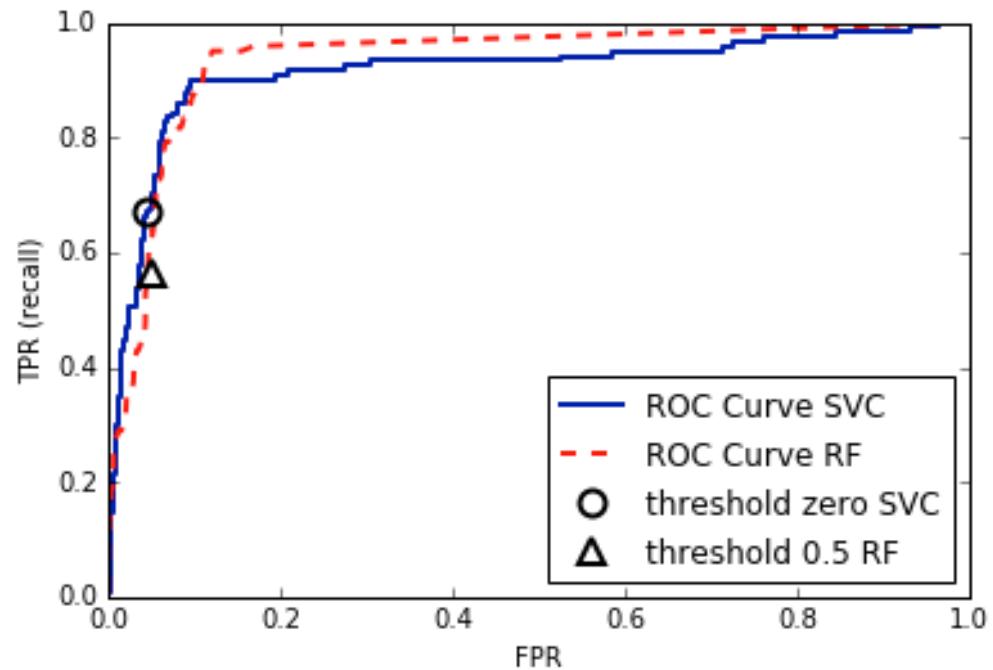


Figure 5-16. Comparing ROC curves for SVM and random forest

Area Under Curve

**Measure of performance
across classification thresholds**

**Notebook
data71200class4.ipynb
[31]**

Decision Thresholds

How much uncertainty the classifier will permit to make a classification

Can be adjusted to change the prioritization between precision and recall

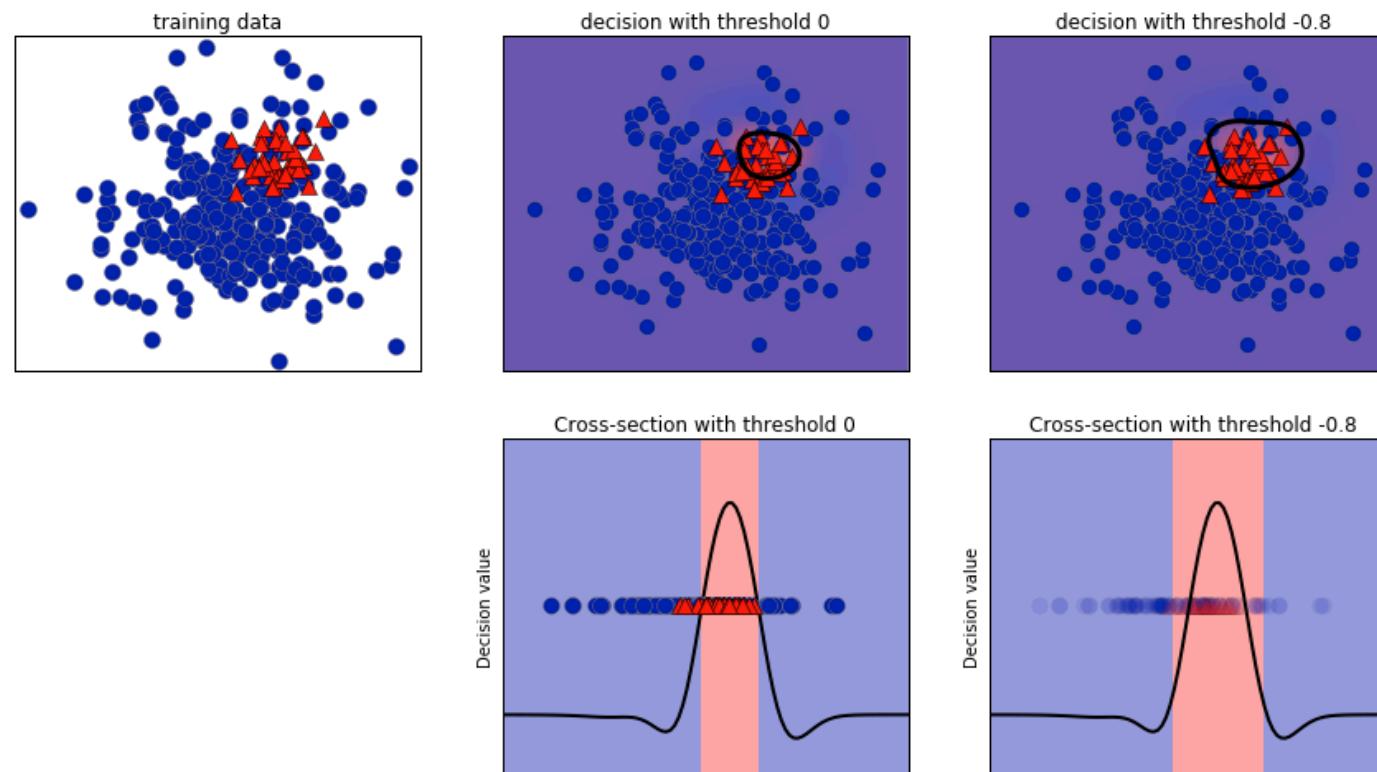


Figure 5-12. Heatmap of the decision function and the impact of changing the decision threshold

Regression Performance Measures

Equation 2-1. Root Mean Square Error (RMSE)

$$\text{RMSE}(\mathbf{X}, h) = \sqrt{\frac{1}{m} \sum_{i=1}^m \left(h(\mathbf{x}^{(i)}) - y^{(i)} \right)^2}$$

m

- m is the number of instances in the dataset you are measuring the RMSE on.
 - For example, if you are evaluating the RMSE on a validation set of 2,000 districts, then $m = 2,000$.

$\mathbf{x}^{(i)}$ $y^{(i)}$

- $\mathbf{x}^{(i)}$ is a vector of all the feature values (excluding the label) of the i^{th} instance in the dataset, and $y^{(i)}$ is its label (the desired output value for that instance).

Regression Performance Measures

$\mathbf{X}^{(i)}$ $y^{(i)}$

- For example, if the first district in the dataset is located at longitude -118.29° , latitude 33.91° , and it has 1,416 inhabitants with a median income of \$38,372, and the median house value is \$156,400 (ignoring the other features for now), then:

$$\mathbf{x}^{(1)} = \begin{pmatrix} -118.29 \\ 33.91 \\ 1,416 \\ 38,372 \end{pmatrix}$$

and:

$$y^{(1)} = 156,400$$

Regression Performance Measures

X

- \mathbf{X} is a matrix containing all the feature values (excluding labels) of all instances in the dataset. There is one row per instance and the i^{th} row is equal to the transpose of $\mathbf{x}^{(i)}$, noted $(\mathbf{x}^{(i)})^T$.⁶
 - For example, if the first district is as just described, then the matrix \mathbf{X} looks like this:

$$\mathbf{X} = \begin{pmatrix} (\mathbf{x}^{(1)})^T \\ (\mathbf{x}^{(2)})^T \\ \vdots \\ (\mathbf{x}^{(1999)})^T \\ (\mathbf{x}^{(2000)})^T \end{pmatrix} = \begin{pmatrix} -118.29 & 33.91 & 1,416 & 38,372 \\ \vdots & \vdots & \vdots & \vdots \end{pmatrix}$$

Regression Performance Measures

h

- h is your system's prediction function, also called a *hypothesis*. When your system is given an instance's feature vector $\mathbf{x}^{(i)}$, it outputs a predicted value $\hat{y}^{(i)} = h(\mathbf{x}^{(i)})$ for that instance (\hat{y} is pronounced "y-hat").
 - For example, if your system predicts that the median housing price in the first district is \$158,400, then $\hat{y}^{(1)} = h(\mathbf{x}^{(1)}) = 158,400$. The prediction error for this district is $\hat{y}^{(1)} - y^{(1)} = 2,000$.

$\text{RMSE}(\mathbf{X}, h)$

- $\text{RMSE}(\mathbf{X}, h)$ is the cost function measured on the set of examples using your hypothesis h .
- Computing the root of a sum of squares (RMSE) corresponds to the *Euclidian norm*: it is the notion of distance you are familiar with. It is also called the ℓ_2 norm, noted $\|\cdot\|_2$ (or just $\|\cdot\|$).

Notebook
data71200class4.ipynb
[32]

Regression Performance Measures

Equation 2-2. Mean Absolute Error

$$\text{MAE}(\mathbf{X}, h) = \frac{1}{m} \sum_{i=1}^m |h(\mathbf{x}^{(i)}) - y^{(i)}|$$

- m
- m is the number of instances in the dataset you are measuring the RMSE on.
- $\mathbf{x}^{(i)}$ $y^{(i)}$
- $\mathbf{x}^{(i)}$ is a vector of all the feature values (excluding the label) of the i^{th} instance in the dataset, and $y^{(i)}$ is its label (the desired output value for that instance).
- Computing the sum of absolutes (MAE) corresponds to the ℓ_1 norm, noted $\|\cdot\|_1$. It is sometimes called the *Manhattan norm* because it measures the distance between two points in a city if you can only travel along orthogonal city blocks.

Notebook
data71200class4.ipynb
[33]

Regression Performance Measures

RMSE is more sensitive to outliers than the MAE

- when outliers are exponentially rare (like in a bell-shaped curve), the RMSE performs very well and is generally preferred

Reading

- Ch 5: “Model Evaluation and Improvement” in
Guido, Sarah and Andreas C. Muller. (2016).
Introduction to Machine Learning with Python,
O'Reilly Media, Inc.

DataCamp for next class

- ▶ *Supervised Learning with scikit-learn*

Videos for next class

- ▶ K-nearest Neighbors: <https://vimeo.com/400660692>
- ▶ Linear Models: <https://vimeo.com/403004687>
- ▶ password: data71200